

Name	xenial (7) time.7.gz
Description	Provided by: manpages_4.04-2_all 
See Also	NAME
Colophon	time - overview of time and timers

DESCRIPTION

Real time and process time

Real time is defined as time measured from some fixed point, either from a standard point in the past (see the description of the Epoch and calendar time below), or from some point (e.g., the start) in the life of a process (elapsed time).

Process time is defined as the amount of CPU time used by a process. This is sometimes divided into user and system components. User CPU time is the time spent executing code in user mode. System CPU time is the time spent by the kernel executing in system mode on behalf of the process (e.g., executing system calls). The `time(1)` command can be used to determine the amount of CPU time consumed during the execution of a program. A program can determine the amount of CPU time it has consumed using `times(2)`, `getrusage(2)`, or `clock(3)`.

The hardware clock

Most computers have a (battery-powered) hardware clock which the kernel reads at boot time in order to initialize the software clock. For further details, see `rtc(4)` and `hwclock(8)`.

The software clock, HZ, and jiffies

The accuracy of various system calls that set timeouts, (e.g., `select(2)`, `sigtimedwait(2)`) and measure CPU time (e.g., `getrusage(2)`) is limited by the resolution of the software clock, a clock maintained by the kernel which measures time in jiffies. The size of a jiffy is determined by the value of the kernel constant `HZ`.

The value of `HZ` varies across kernel versions and hardware platforms. On i386 the situation is as follows: on kernels up to and including 2.4.x, `HZ` was 100, giving a jiffy value of 0.01 seconds; starting with 2.6.0, `HZ` was raised to 1000, giving a jiffy of 0.001 seconds. Since kernel 2.6.13, the `HZ` value is a kernel configuration parameter and can be 100, 250 (the default) or 1000, yielding a jiffies value of, respectively, 0.01, 0.004, or 0.001 seconds. Since kernel 2.6.20, a further frequency is available: 300, a number that divides evenly for the common video frame rates (PAL, 25 HZ; NTSC, 30 HZ).

The `times(2)` system call is a special case. It reports times with a granularity defined by the kernel constant `USER_HZ`. User-space applications can determine the value of this constant using `sysconf(SC_CLK_TCK)`.

High-resolution timers

Before Linux 2.6.21, the accuracy of timer and sleep system calls

(see below) was also limited by the size of the jiffy.

Since Linux 2.6.21, Linux supports high-resolution timers (HRTs), optionally configurable via `CONFIG_HIGH_RES_TIMERS`. On a system that supports HRTs, the accuracy of sleep and timer system calls is no longer constrained by the jiffy, but instead can be as accurate as the hardware allows (microsecond accuracy is typical of modern hardware). You can determine whether high-resolution timers are supported by checking the resolution returned by a call to `clock_getres(2)` or looking at the "resolution" entries in [/proc/timer_list](#).

HRTs are not supported on all hardware architectures. (Support is provided on x86, arm, and powerpc, among others.)

The Epoch

UNIX systems represent time in seconds since the Epoch, 1970-01-01 00:00:00 +0000 (UTC).

A program can determine the calendar time using `gettimeofday(2)`, which returns time (in seconds and microseconds) that have elapsed since the Epoch; `time(2)` provides similar information, but only with accuracy to the nearest second. The system time can be changed using `settimeofday(2)`.

Broken-down time

Certain library functions use a structure of type `tm` to represent broken-down time, which stores time value separated out into distinct components (year, month, day, hour, minute, second, etc.). This structure is described in `ctime(3)`, which also describes functions that convert between calendar time and broken-down time. Functions for converting between broken-down time and printable string representations of the time are described in `ctime(3)`, `strptime(3)`, and `strptime(3)`.

Sleeping and setting timers

Various system calls and functions allow a program to sleep (suspend execution) for a specified period of time; see `nanosleep(2)`, `clock_nanosleep(2)`, and `sleep(3)`.

Various system calls allow a process to set a timer that expires at some point in the future, and optionally at repeated intervals; see `alarm(2)`, `getitimer(2)`, `timerfd_create(2)`, and `timer_create(2)`.

Timer slack

Since Linux 2.6.28, it is possible to control the "timer slack" value for a thread. The timer slack is the length of time by which the kernel may delay the wake-up of certain system calls that block with a timeout. Permitting this delay allows the kernel to coalesce wake-up events, thus possibly reducing the number of system wake-ups and saving power. For more details, see the description of `PR_SET_TIMERSLACK` in `prctl(2)`.

SEE ALSO

`date(1)`, `time(1)`, `adjtimex(2)`, `alarm(2)`, `clock_gettime(2)`, `clock_nanosleep(2)`, `getitimer(2)`, `getrlimit(2)`, `getrusage(2)`, `gettimeofday(2)`, `nanosleep(2)`, `stat(2)`, `time(2)`, `timer_create(2)`, `timerfd_create(2)`, `times(2)`, `utime(2)`, `adjtime(3)`,

```
clock(3),
clock_getcpuclockid(3), ctime(3), pthread_getcpuclockid(3),
sleep(3), strftime(3),
strptime(3), timeradd(3), usleep(3), rtc(4), hwclock(8)
```

COLOPHON

This page is part of release 4.04 of the Linux [man-pages](#) project. A description of the project, information about reporting bugs, and the latest version of this page, can be found at <http://www.kernel.org/doc/man-pages/>.

